

Neural Turing Machines

and memory-augmented techniques

Ananth Mahadevan, Christabella Irwanto

Agenda

all about Neural Turing Machines
(NTMs)

- Motivation
- How NTMs work
- Experiments
- Discussion

— — —

Motivation


- Neuroscience angle: "working memory" is limited
 - Cognitive system with limited capacity for **temporarily** storing and manipulating information
 - Short term memory cannot manipulate the information
- For a hard mathematical problem/memorizing many numbers, ~~we~~ most of us (except Ananth) would need to **offload some computation to pen and paper**
- Like computer programs
- Learn how to manage memory to solve a problem: sort of "learning to program"
 - NTM learns its own basic algorithms for tasks such as copying, sorting, and associative recall

NTM

- Entire structure is differentiable
- Part of broader trend to "**differentiable programming**"
 - incorporation of **classic data structures** (e.g. RAM, stacks, queues) into **gradient-based learning systems**
 - Increasing memory capacity in NTM is easier than in pure LSTM networks.
- Possibility of combining the best of program induction and deep learning
 - Structured representations – *objects, forces, agents, causality, compositionality* – help explain important facets of human learning and thinking
 - Deep learning systems not shown to work with these representations, but shown surprising effectiveness of gradient descent in large models

Neural Turing Machine (NTM)

- Neural network *controller* (e.g. feed forward or RNN)
 - also outputs "heads" that parameterize reads and writes to/from the...
- Memory bank! (stores processed information)



NTM unfolded

- Controller accepts input x and read vector r , outputs read and write heads

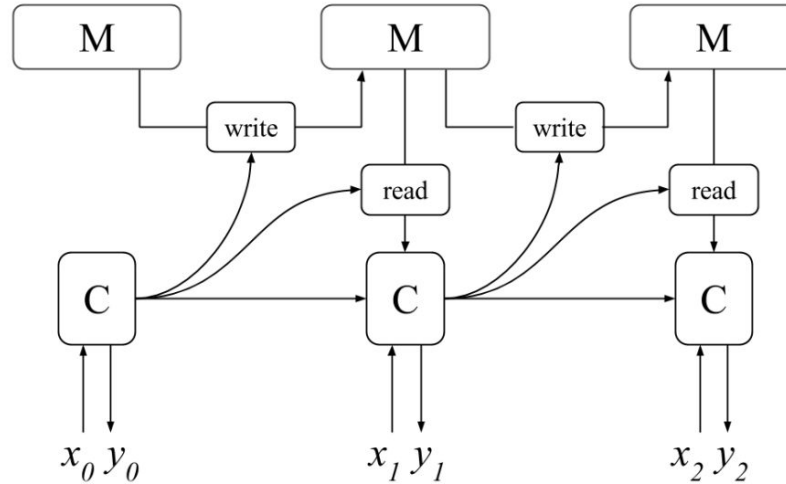


FIGURE 2.13: The neural Turing machine unrolled through time (Olah and Carter, 2016).

Reading

- At time \mathbf{t} , normalized weight vector \mathbf{w}_t controls how much attention to give to different memory locations

$$\sum_i w_t(i) = 1, \quad 0 \leq w_t(i) \leq 1, \quad \forall i.$$

- Read vector \mathbf{r}_t is simply a sum weighted by attention intensity

$$\mathbf{r}_t \leftarrow \sum_i w_t(i) \mathbf{M}_t(i).$$

- Differentiable with respect to both the weight and the memory bank
- A form of attention implemented as memory addressing

Writing

- Inspired by input and forget gates in LSTM
- Given weight vector \mathbf{w}_t emitted by a write head at time t , we
 - **erase** with \mathbf{e}_t

$$\tilde{\mathbf{M}}_t(i) \leftarrow \mathbf{M}_{t-1}(i) [1 - w_t(i)\mathbf{e}_t]$$

- **add** with \mathbf{a}_t


$$\mathbf{M}_t(i) \leftarrow \tilde{\mathbf{M}}_t(i) + w_t(i) \mathbf{a}_t$$

Addressing mechanisms

- How is \mathbf{w}_t produced?
- \mathbf{w}_t is updated through a series of four intermediate smooth operations
 - *content addressing*
 - *interpolation*
 - *convolutional shift*
 - *sharpening*
- Operations depend on parameters from the controller
 - functions of hidden state \mathbf{h}_t emitted by the controller.

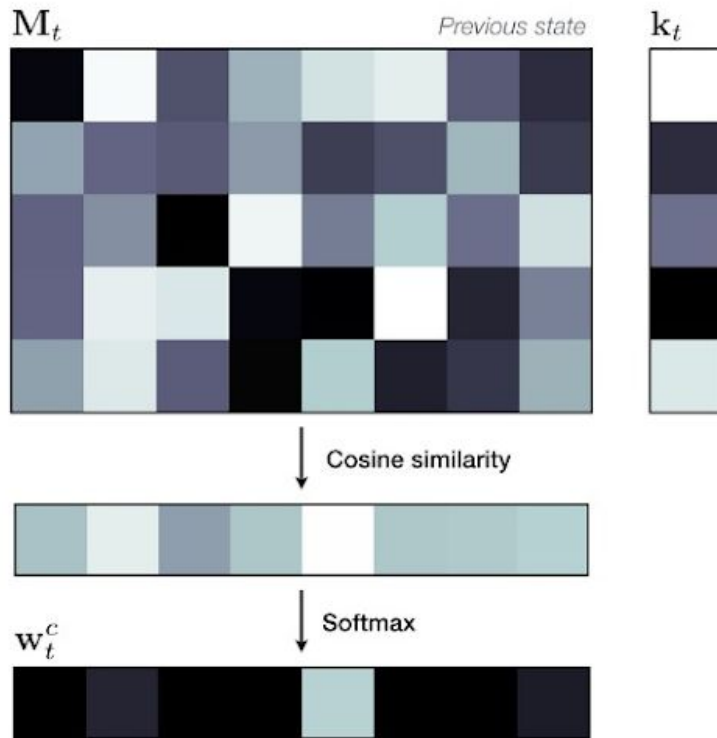
Parameters for weight (\mathbf{w}_t) updates

- 5 parameters specific to each read/write head
- Each box is a 1-layer neural network with some activation function



Content addressing

$$w_t^c(i) \leftarrow \text{softmax} \left(\beta_t \cdot \frac{\mathbf{k}_t \cdot \mathbf{M}_t(i)}{\|\mathbf{k}_t\| \cdot \|\mathbf{M}_t(i)\|} \right)$$




- retrieve specific informations in memory
- compute cosine similarity between key vector \mathbf{k}_t extracted by controller from input and memory
- then normalized by softmax
 - with strength multiplier β_t to amplify or attenuate the focus of the distribution

— — —

Interpolation

$$\mathbf{w}_t^g \leftarrow g_t \mathbf{w}_t^c + (1 - g_t) \mathbf{w}_{t-1}$$




- interpolation gate scalar \mathbf{g}_t blends newly-generated content-based weight vector \mathbf{w}_t^c with weight vector from previous time step \mathbf{w}_{t-1}

— — —

Convolutional shift


$$\tilde{w}_t(i) \leftarrow \sum_j w_t^g(j) \cdot s_t(i - j)$$



- location-based addressing done by 1-d convolution of \mathbf{w}_t^g with kernel $\mathbf{s}_t(\cdot)$, a function of the position offset $\mathbf{i} - \mathbf{j}$

Sharpening

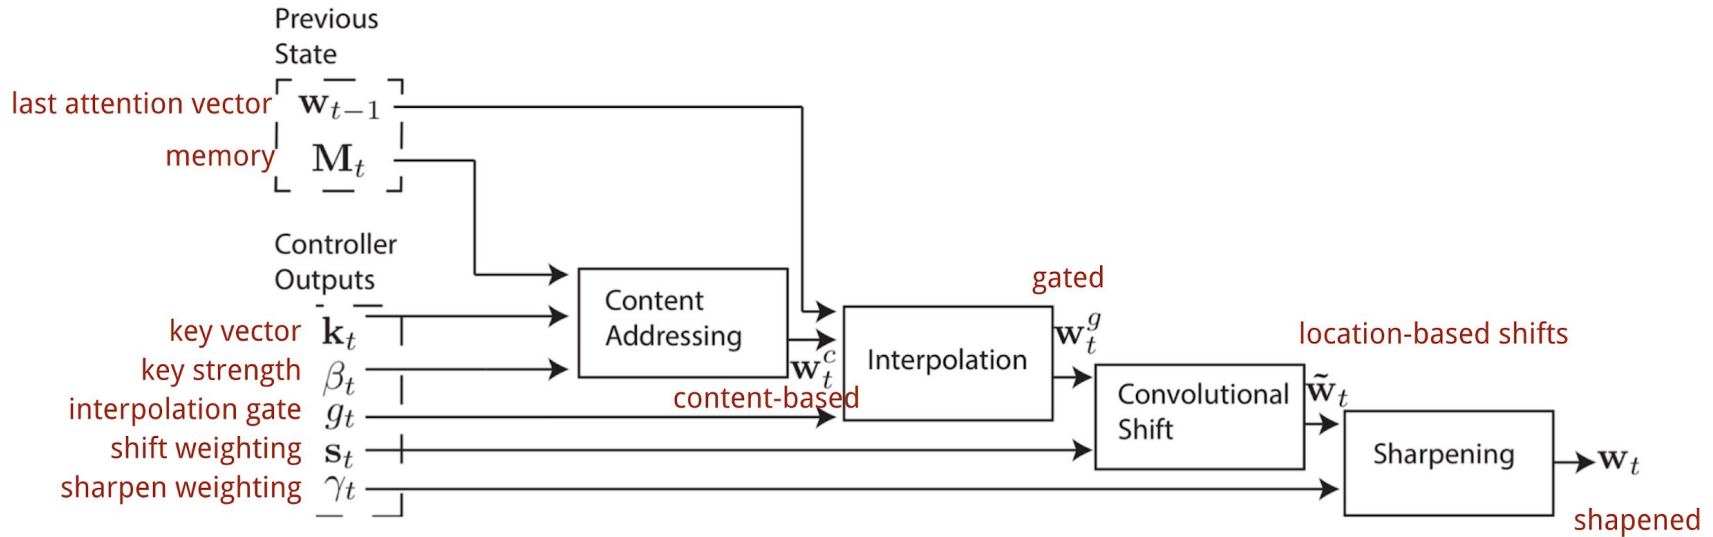
$$w_t(i) \propto \tilde{w}_t(i)^{\gamma_t} = \frac{\tilde{w}_t(i)^{\gamma_t}}{\sum_{j=1}^N \tilde{w}_t(j)^{\gamma_t}}$$



- shifted attention vector is sharpened with a sharpening scalar $\gamma_t \geq 1$

Flow diagram of addressing mechanisms

- Complete process of generating the attention vector \mathbf{w}_t



Controller Network Architecture

Feed forward vs recurrent:


- LSTM has its own internal memory complementary to **M**
- Feed forward has better transparency

Analogy:

- Controller \Leftrightarrow CPU
- Memory Matrix \Leftrightarrow RAM
- Hidden Activations of LSTM \Leftrightarrow Registers in Processor

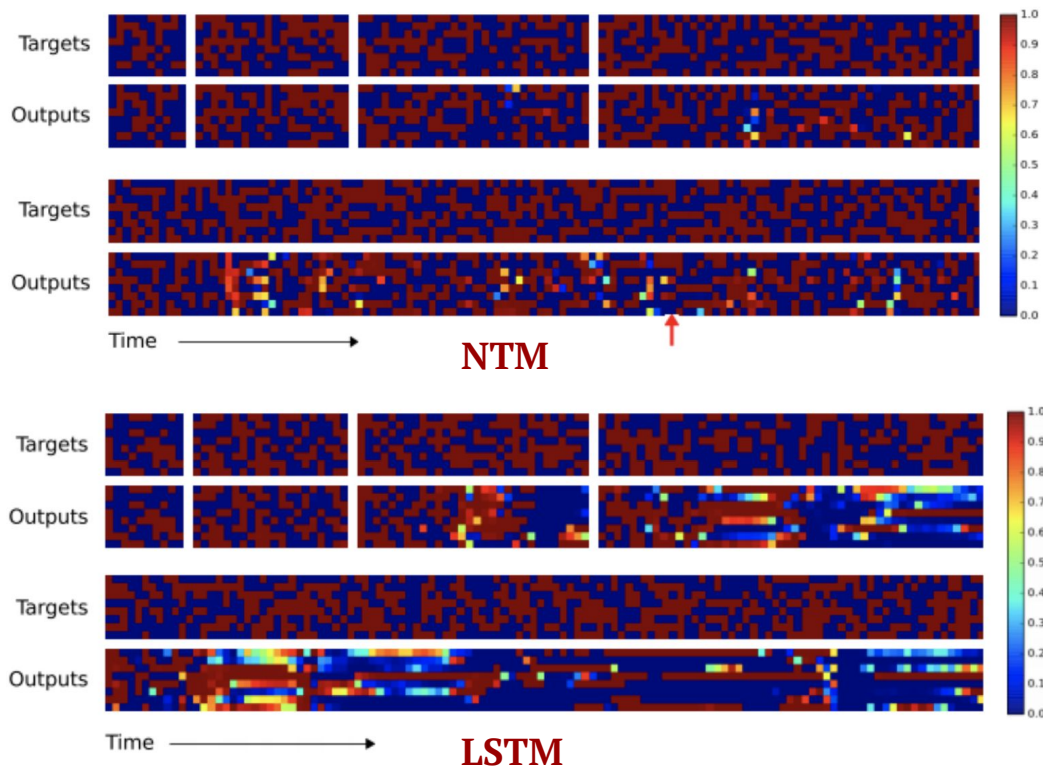
Experiments: Copy

- RNNs have struggled to remember information over long periods
- Given n random eight bit vectors followed by a delimiter flag, repeat it
- "Cost per sequence": number of bits incorrectly recalled over a sequence



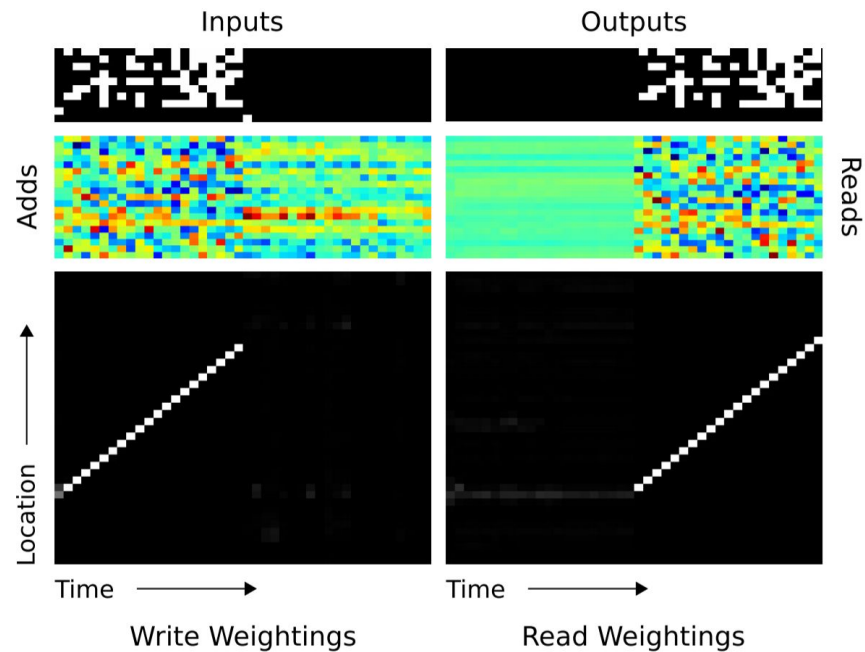
Experiments: Copy

- **Training**
sequences of
lengths 1 to 20
- **Test** with lengths
10, 20, 30, 50, 120
- [blue, red] -> [0, 1]
- NTM far fewer
errors on longer
sequences; scales
better




Experiments: Copy

- Visualize how the NTM reads from and writes to \mathbf{M}



Experiments: Repeat Copy

- Like a **nested for loop**: repeatedly **copy** sequence some **x number of times**
- Training input: random-length sequence of 8-bit binary vectors + scalar x from 1-10



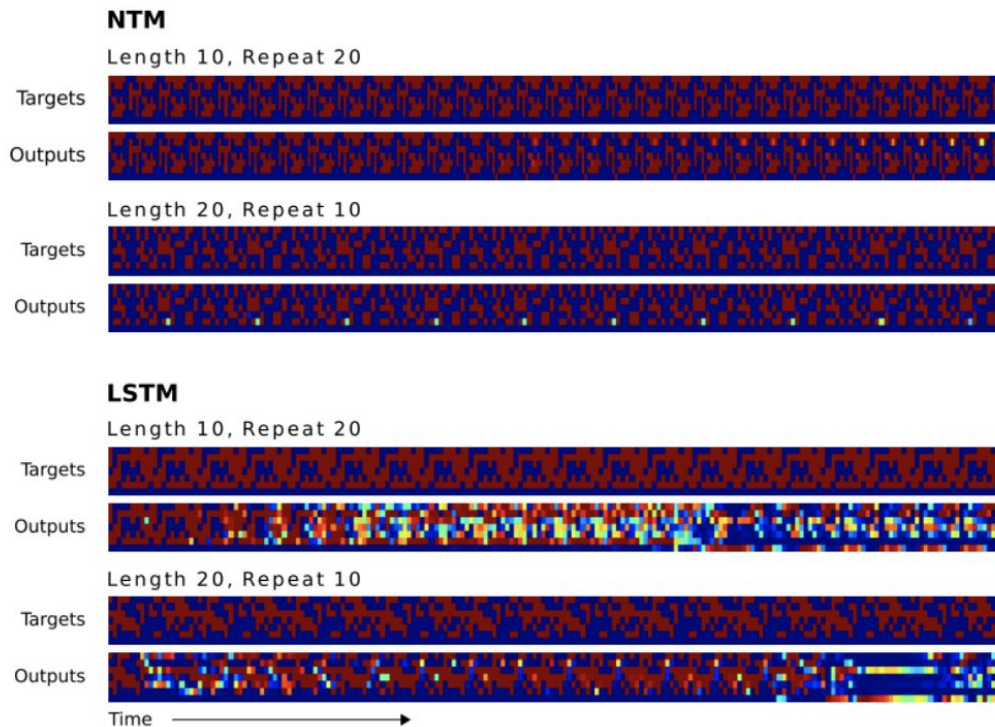
Experiments: Repeat Copy

Increased # repeats:

- NTM much better, but still keeps falsely predicting end of sequence (emits delimiter flag after every repetition beyond the 11th)

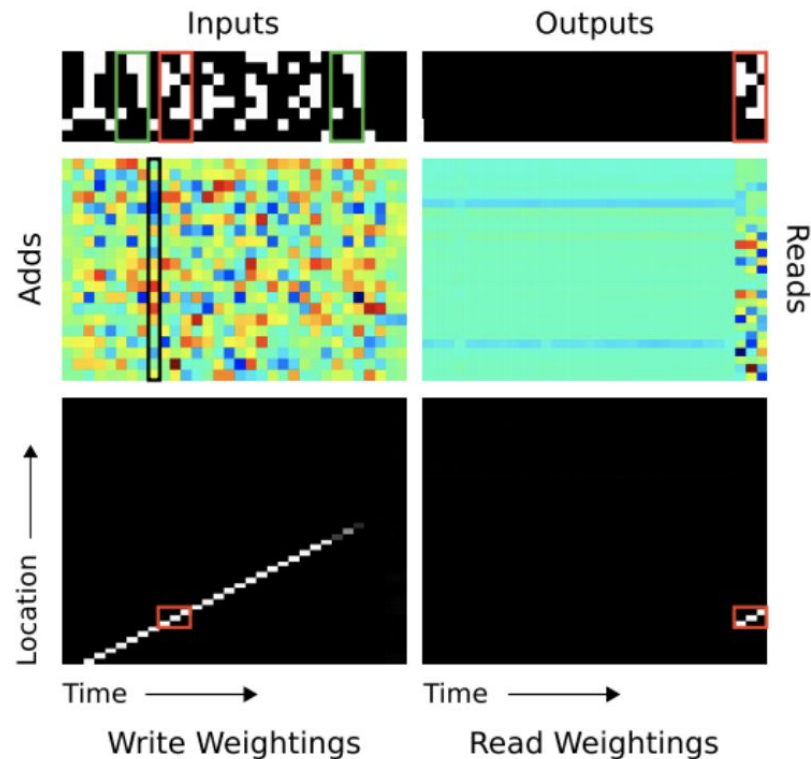
Increased length:

- NTM generalizes much better



Experiment: Associative recall

- Can NTMs learn "indirection" i.e. one data item pointing to another?
- **input:**
 - list of items
 - query of one of the items (green)
- **output:**
 - next item in the list (red)
- at delimiter flag, NTM forms compressed representation (black box in "Adds") of each item



Experiment: Associative recall

- **Feed forward-controller NTM outperforms LSTM-controller NTM**
 - suggests that NTM's memory is a superior data storage system than LSTM's internal state




Figure 10: Associative Recall Learning Curves for NTM and LSTM.

Segue to code

Experiment: Dynamic n-grams


- Whether NTMs could learn posterior predictive distributions
- N-grams (sequences of N items), which given previous items in the sequence, determine some probability distribution over the next item in the sequence
- Optimal estimator:

$$P(B = 1|N_1, N_0, \mathbf{c}) = \frac{N_1 + \frac{1}{2}}{N_1 + N_0 + 1}$$

- **B** is the next bit, **c** is the previous 5-bit context, and **N0** and **N1** are the number of zeros and ones observed after **c**
- NTM achieves small, but significant performance advantage over LSTM, but never quite reaches optimum cost


Experiment: Dynamic n-grams

- Top row: test sequence
- Rows below: predictive distributions by optimal estimator, NTM, and LSTM
- NTM predictions mostly indistinguishable from optimal (although some clear mistakes e.g. red arrows)
- LSTM good but appears to diverge further as sequence progresses
 - speculate that LSTM "forgets" observations at start



Experiment: Dynamic n-grams

- red and green arrows where same context is observed (“00010” for green, “01111” for red)
- same location read then written to
- Network uses writes to keep count of the fraction of ones and zeros following each context in the sequence so far
- Add vectors clearly anti-correlated at 0s and 1s, suggesting a distributed “counter.”



Experiment: Priority sort

Whether the NTM can sort data—an important elementary algorithm





Figure 16: Example Input and Target Sequence for the Priority Sort Task. The input sequence contains random binary vectors and random scalar priorities. The target sequence is a subset of the input vectors sorted by the priorities.

Experiment: Priority sort

Again, NTMs > LSTM



Implementational Checkpoint

- Took 4 years to find a stable implementation
- Paper just on implementing NTMs

arXiv.org > cs > arXiv:1807.08518

Search...

Help | Adv

Computer Science > Machine Learning

Implementing Neural Turing Machines

[Mark Collier](#), [Joeran Beel](#)


(Submitted on 23 Jul 2018 (v1), last revised 26 Jul 2018 (this version, v3))

Neural Turing Machines (NTMs) are an instance of Memory Augmented Neural Networks, a new class of recurrent neural networks which decouple computation from memory by introducing an external memory unit. NTMs have demonstrated superior performance over Long Short-Term Memory Cells in several sequence learning tasks. A number of open source implementations of NTMs exist but are unstable during training and/or fail to replicate the reported performance of NTMs. This paper presents the details of our successful implementation of a NTM. Our implementation learns to solve three sequential learning tasks from the original NTM paper. We find that the choice of memory contents initialization scheme is crucial in successfully implementing a NTM. Networks with memory contents initialized to small constant values converge on average 2 times faster than the next best memory contents initialization scheme.

Differentiable Neural Computer (DNC)

- Successor of NTMs; more generalizable network structure
- Wide range of tasks including natural language understanding, basic inference, planning, etc.
 - E.g. Solving block puzzles
 - E.g. Finding paths between nodes in a graph after memorizing the graph
- Has ability to **free** and **allocate** memory, i.e. to reuse memory locations
- Using **temporal links**, can search by time of writes

DNC



References

- <https://medium.com/snips-ai/ntm-lasagne-a-library-for-neural-turing-machines-in-lasagne-2cdce6837315>
- <https://github.com/MarkPKCollier/NeuralTuringMachine>
- <https://github.com/loudinthecloud/pytorch-ntm>
- <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>
- <http://www.robots.ox.ac.uk/~tvq/publications/talks/NeuralTuringMachines.pdf>
- https://rylanschaeffer.github.io/content/research/neural_turing_machine/main.html
- Lake, B., Ullman, T., Tenenbaum, J., & Gershman, S. (2017). Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40, E253. doi:10.1017/S0140525X16001837
- Graves, Alex; Wayne, Greg; Danihelka, Ivo (2014). "Neural Turing Machines". [arXiv:1410.5401](#)
- Graves, Alex; Wayne, Greg; Reynolds, Malcolm; Harley, Tim; Danihelka, Ivo; Grabska-Barwińska, Agnieszka; Colmenarejo, Sergio Gómez; Grefenstette, Edward; Ramalho, Tiago (2016-10-12). "[Hybrid computing using a neural network with dynamic external memory](#)". *Nature*. **538** (7626): 471–476. [Bibcode:2016Natur.538..471G](#). [doi:10.1038/nature20101](#). [ISSN 1476-4687](#). [PMID 27732574](#).